

dl-ex-5

May 1, 2023

```
[ ]: from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.models import Model
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
```

```
[2]: # Load the CIFAR-10 dataset
(X_train, y_train), (X_test, y_test) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 3s 0us/step

```
[3]: # Normalize the pixel values
X_train = X_train / 255.0
X_test = X_test / 255.0
```

```
[4]: # Convert the labels to one-hot encoded vectors
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)
```

```
[5]: # Load the pre-trained VGG-16 model
vgg16 = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58889256/58889256 [=====] - 0s 0us/step

```
[6]: # Freeze the weights of the pre-trained layers
for layer in vgg16.layers:
    layer.trainable = False
```

```
[7]: # Add new layers for classification
x = Flatten()(vgg16.output)
x = Dense(units=256, activation='relu')(x)
x = Dense(units=10, activation='softmax')(x)
```

```
[8]: # Create a new model that combines the pre-trained VGG-16 model with the new  
classification layers  
model = Model(inputs=vgg16.input, outputs=x)
```

```
[9]: # Compile the model  
model.compile(optimizer='adam', loss='categorical_crossentropy',  
metrics=['accuracy'])
```

```
[10]: # Train the model  
model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test,  
y_test))
```

```
Epoch 1/10  
1563/1563 [=====] - 772s 493ms/step - loss: 1.3405 -  
accuracy: 0.5327 - val_loss: 1.2420 - val_accuracy: 0.5626  
Epoch 2/10  
1563/1563 [=====] - 766s 490ms/step - loss: 1.1675 -  
accuracy: 0.5912 - val_loss: 1.1956 - val_accuracy: 0.5815  
Epoch 3/10  
1563/1563 [=====] - 765s 489ms/step - loss: 1.1027 -  
accuracy: 0.6124 - val_loss: 1.1382 - val_accuracy: 0.5961  
Epoch 4/10  
1563/1563 [=====] - 766s 490ms/step - loss: 1.0519 -  
accuracy: 0.6333 - val_loss: 1.1416 - val_accuracy: 0.6017  
Epoch 5/10  
1563/1563 [=====] - 764s 489ms/step - loss: 1.0087 -  
accuracy: 0.6465 - val_loss: 1.1113 - val_accuracy: 0.6127  
Epoch 6/10  
1563/1563 [=====] - 763s 488ms/step - loss: 0.9701 -  
accuracy: 0.6597 - val_loss: 1.1203 - val_accuracy: 0.6080  
Epoch 7/10  
1563/1563 [=====] - 766s 490ms/step - loss: 0.9373 -  
accuracy: 0.6703 - val_loss: 1.1163 - val_accuracy: 0.6121  
Epoch 8/10  
1563/1563 [=====] - 766s 490ms/step - loss: 0.9037 -  
accuracy: 0.6813 - val_loss: 1.1054 - val_accuracy: 0.6140  
Epoch 9/10  
1563/1563 [=====] - 764s 489ms/step - loss: 0.8730 -  
accuracy: 0.6916 - val_loss: 1.1292 - val_accuracy: 0.6104  
Epoch 10/10  
1563/1563 [=====] - 765s 489ms/step - loss: 0.8410 -  
accuracy: 0.7051 - val_loss: 1.1214 - val_accuracy: 0.6162
```

```
[10]: <keras.callbacks.History at 0x7f4ab179fd60>
```

```
[11]: from tensorflow.keras.applications.vgg19 import VGG19  
from tensorflow.keras.layers import Dense, Flatten
```

```

from tensorflow.keras.models import Model
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical

[12]: # Load the CIFAR-10 dataset
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

[13]: # Normalize the pixel values
X_train = X_train / 255.0
X_test = X_test / 255.0

[14]: # Convert the labels to one-hot encoded vectors
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)

[15]: # Load the pre-trained VGG-19 model
vgg19 = VGG19(weights='imagenet', include_top=False, input_shape=(32, 32, 3))

Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5
80134624/80134624 [=====] - 1s 0us/step

[16]: # Freeze the weights of the pre-trained layers
for layer in vgg19.layers:
    layer.trainable = False

[17]: # Add new layers for classification
x = Flatten()(vgg19.output)
x = Dense(units=256, activation='relu')(x)
x = Dense(units=10, activation='softmax')(x)

[18]: # Create a new model that combines the pre-trained VGG-19 model with the new
      ↴classification layers
model = Model(inputs=vgg19.input, outputs=x)

[19]: # Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy',
      ↴metrics=['accuracy'])

[20]: # Train the model
model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test,
      ↴y_test))

```

Epoch 1/10
 1563/1563 [=====] - 940s 601ms/step - loss: 1.3967 -
 accuracy: 0.5070 - val_loss: 1.2662 - val_accuracy: 0.5582
 Epoch 2/10

```
1563/1563 [=====] - 935s 598ms/step - loss: 1.2170 -  
accuracy: 0.5725 - val_loss: 1.2204 - val_accuracy: 0.5724  
Epoch 3/10  
1563/1563 [=====] - 983s 629ms/step - loss: 1.1562 -  
accuracy: 0.5945 - val_loss: 1.1994 - val_accuracy: 0.5796  
Epoch 4/10  
1563/1563 [=====] - 984s 630ms/step - loss: 1.1091 -  
accuracy: 0.6081 - val_loss: 1.1686 - val_accuracy: 0.5929  
Epoch 5/10  
1563/1563 [=====] - 984s 630ms/step - loss: 1.0703 -  
accuracy: 0.6229 - val_loss: 1.1450 - val_accuracy: 0.6033  
Epoch 6/10  
1563/1563 [=====] - 984s 630ms/step - loss: 1.0366 -  
accuracy: 0.6344 - val_loss: 1.1701 - val_accuracy: 0.5942  
Epoch 7/10  
1563/1563 [=====] - 984s 629ms/step - loss: 1.0047 -  
accuracy: 0.6461 - val_loss: 1.1477 - val_accuracy: 0.6057  
Epoch 8/10  
1563/1563 [=====] - 932s 597ms/step - loss: 0.9770 -  
accuracy: 0.6550 - val_loss: 1.1215 - val_accuracy: 0.6171  
Epoch 9/10  
1563/1563 [=====] - 981s 628ms/step - loss: 0.9480 -  
accuracy: 0.6659 - val_loss: 1.1326 - val_accuracy: 0.6099  
Epoch 10/10  
1563/1563 [=====] - 937s 599ms/step - loss: 0.9221 -  
accuracy: 0.6738 - val_loss: 1.1427 - val_accuracy: 0.6098
```

[20]: <keras.callbacks.History at 0x7f4aad716140>