

dl-ex-6

May 1, 2023

```
[12]: import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np

[13]: # Load the CIFAR-10 dataset
(X_train, y_train), (X_test, y_test) = tf.keras.datasets.cifar10.load_data()

[14]: # Normalize the pixel values
X_train = X_train / 255.0
X_test = X_test / 255.0

[15]: # Convert the labels to one-hot encoded vectors
y_train = tf.keras.utils.to_categorical(y_train, num_classes=10)
y_test = tf.keras.utils.to_categorical(y_test, num_classes=10)

[16]: # Define the CNN model with different activation functions
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(filters=32, kernel_size=(3,3), activation='relu', ↴
    ↪input_shape=(32, 32, 3)),
    tf.keras.layers.MaxPooling2D(pool_size=(2,2)),
    tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3), activation='sigmoid'),
    tf.keras.layers.MaxPooling2D(pool_size=(2,2)),
    tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3), activation='tanh'),
    tf.keras.layers.MaxPooling2D(pool_size=(2,2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(units=256, activation='relu'),
    tf.keras.layers.Dense(units=10, activation='softmax')])

[17]: # Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', ↴
    ↪metrics=['accuracy'])

[18]: # Train the model
model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, ↴
    ↪y_test))
```

Epoch 1/10
1563/1563 [=====] - 99s 62ms/step - loss: 1.7032 -

```
accuracy: 0.3759 - val_loss: 1.4279 - val_accuracy: 0.4778
Epoch 2/10
1563/1563 [=====] - 97s 62ms/step - loss: 1.3110 -
accuracy: 0.5293 - val_loss: 1.2034 - val_accuracy: 0.5704
Epoch 3/10
1563/1563 [=====] - 94s 60ms/step - loss: 1.1495 -
accuracy: 0.5916 - val_loss: 1.1504 - val_accuracy: 0.5931
Epoch 4/10
1563/1563 [=====] - 92s 59ms/step - loss: 1.0157 -
accuracy: 0.6406 - val_loss: 1.0547 - val_accuracy: 0.6290
Epoch 5/10
1563/1563 [=====] - 93s 59ms/step - loss: 0.9076 -
accuracy: 0.6797 - val_loss: 0.9732 - val_accuracy: 0.6563
Epoch 6/10
1563/1563 [=====] - 94s 60ms/step - loss: 0.8136 -
accuracy: 0.7149 - val_loss: 0.9637 - val_accuracy: 0.6690
Epoch 7/10
1563/1563 [=====] - 92s 59ms/step - loss: 0.7302 -
accuracy: 0.7442 - val_loss: 0.9489 - val_accuracy: 0.6756
Epoch 8/10
1563/1563 [=====] - 97s 62ms/step - loss: 0.6504 -
accuracy: 0.7720 - val_loss: 0.9490 - val_accuracy: 0.6815
Epoch 9/10
1563/1563 [=====] - 97s 62ms/step - loss: 0.5826 -
accuracy: 0.7956 - val_loss: 0.9572 - val_accuracy: 0.6952
Epoch 10/10
1563/1563 [=====] - 99s 64ms/step - loss: 0.5128 -
accuracy: 0.8190 - val_loss: 0.9552 - val_accuracy: 0.7023
```

[18]: <keras.callbacks.History at 0x7f6c28fe1e70>

[19]: # Extract the activations of the convolutional layers
activation_model = tf.keras.models.Model(inputs=model.input, outputs=[layer.
↳output for layer in model.layers if isinstance(layer, tf.keras.layers.
↳Conv2D)])

[20]: # Choose an image from the test set to visualize
image = X_test[0]
image = np.expand_dims(image, axis=0)

[21]: # Generate heatmaps for each activation function
activations = activation_model.predict(image)
relu_activation = activations[0]
sigmoid_activation = activations[1]
tanh_activation = activations[2]

```
1/1 [=====] - 0s 103ms/step
```

```
[22]: # Plot the heatmaps for each activation function
fig, axs = plt.subplots(nrows=1, ncols=3, figsize=(10, 5))
axs[0].imshow(relu_activation[0,:,:,:], cmap='jet')
axs[0].set_title('ReLU activation')
axs[1].imshow(sigmoid_activation[0,:,:,:], cmap='jet')
axs[1].set_title('Sigmoid activation')
axs[2].imshow(tanh_activation[0,:,:,:], cmap='jet')
axs[2].set_title('Tanh activation')
plt.show()
```

